

Order statistics: from simple to $O(n)$

Banin, M.

January 23, 2018

When in doubt

This does what you think it does in expected $O(n)$.

```
int get_element(const vector<int>& vec, size_t pos){
    assert(pos < vec.size());
    vector<int> myvector(vec.begin(), vec.end());
    std::nth_element (myvector.begin(),
                      myvector.begin() + pos,
                      myvector.end());
    return myvector[pos];
}
```

How to do it by hand

```
template <class Iter, class T>
void my_element(Iter first, Iter nth, Iter last) {
    while (last - first > 3) {
        T pivot = *(first + (last - first) / 2);

        Iter ll = partition(first, last,
            [pivot](const T& a) -> bool{ return a < pivot; });
        Iter rr = partition(ll, last,
            [pivot](const T&a) -> bool{ return a == pivot; });

        if (nth < ll) last = ll;
        else if(nth >= ll && nth < rr) return;
        else if(nth >= rr) first = rr;
    }
    std::sort(first, last);
}
```


What STL does

```
template<typename _RandomAccessIterator, typename _Size, typename _Compare>
void
__introsortselect(_RandomAccessIterator __first, _RandomAccessIterator __nth,
                 _RandomAccessIterator __last, _Size __depth_limit,
                 _Compare __comp)
{
    while (__last - __first > 3)
    {
        if (__depth_limit == 0)
        {
            std::__heap_select(__first, __nth + 1, __last, __comp);
            // Place the nth largest element in its final position.
            std::iter_swap(__first, __nth);
            return;
        }
        --__depth_limit;
        _RandomAccessIterator __cut =
            std::__unguarded_partition_pivot(__first, __last, __comp);
        if (__cut <= __nth)
            __first = __cut;
        else
            __last = __cut;
    }
    std::__insertion_sort(__first, __last, __comp);
}
```

```
/// This is a helper function...
```

```
template<typename _RandomAccessIterator, typename _Compare>
    _RandomAccessIterator
    __unguarded_partition(_RandomAccessIterator __first,
                        _RandomAccessIterator __last,
                        _RandomAccessIterator __pivot, _Compare __comp)
{
    while (true)
    {
        while (__comp(__first, __pivot))
            ++__first;
        --__last;
        while (__comp(__pivot, __last))
            --__last;
        if (!(__first < __last))
            return __first;
        std::iter_swap(__first, __last);
        ++__first;
    }
}
```

Median-of medians (Blum, Floyd, Pratt, Rivest, Tarjan)



A bit of complexity theory

- ▶ $\min x = 5^i : x > n = O(n)$
- ▶ $\frac{2}{10} + \frac{7}{10} < 1$

Makeshift strict $O(n)$ implementation

```
template <class Iter, class T>
void lintime_element(Iter first, Iter nth, Iter last) {
    while (last - first > 5) {
        std::vector<T> medians;
        for(Iter it = first; it < last; it += 5){
            sort(it, std::min(it + 5, last));
            medians.push_back(*std::min(it + 2, last - 1));
        }

        T padding = *max_element(medians.begin(), medians.end());
        while(medians.size() % 5 != 0)
            medians.push_back(padding);

        Iter mid = medians.begin() + (medians.end() - medians.begin()) / 2;
        lintime_element<typename std::vector<T>::iterator, T>(medians.begin(),
                                                                mid,
                                                                medians.end());

        T pivot = *mid;

        ...
    }
}
```


Benchmark

On 256K-element arrays, we see:

Implementation	Average runtime
sorting in $O(n \log n)$	52 μsec
makeshift strict $O(n)$	34 μsec
makeshift expected $O(n)$	17 μsec
nth_element call	15 μsec
generating the array	11 μsec

References

- ▶ http://www.cplusplus.com/reference/algorithm/nth_element/
- ▶ <https://en.wikipedia.org/wiki/Quickselect>
- ▶ Blum, M.; Floyd, R. W.; Pratt, V. R.; Rivest, R. L.; Tarjan, R. E. (August 1973). "Time bounds for selection"
- ▶ Noriyuki Kurosawa. (2016) Quicksort with median of medians considered practic